

CS441: Applied ML - HW 3

Part 1: CLIP: Contrastive Language-Image Pretraining

Include all the code for Part 1 in this section

1.1 Prepare data

[Here](#) is the json file you need for labels of flowers 102

```
import json
import os
import os.path as osp
import numpy as np
from google.colab import drive
import torch
from torchvision.datasets import Flowers102
%matplotlib inline
from matplotlib import pyplot as plt

drive.mount('/content/drive')
datadir = "/content/drive/My Drive/CS441/hw3/"

Mounted at /content/drive

def load_flower_data(img_transform=None):
    drive.mount('/content/drive')
    datadir = "/content/drive/My Drive/CS441/hw3/" #/content/drive/My Drive/CS441/hw3/"
    if os.path.isdir(datadir+ "flowers-102"):
        do_download = False
    else:
        do_download = True
    train_set = Flowers102(root=datadir, split='train', transform=img_transform, download=do_download)
    test_set = Flowers102(root=datadir, split='val', transform=img_transform, download=do_download)
    classes = json.load(open(osp.join("/content/drive/My Drive/CS441/hw3/", "flowers102_classes.json")))

    return train_set, test_set, classes

# READ ME! This takes some time (a few minutes), so if you are using Colabs,
# first set to use GPU: Edit->Notebook Settings->Hardware Accelerator=GPU, and restart instance

# Data structure details
# flower_train[n][0] is the nth train image
# flower_train[n][1] is the nth train label
# flower_test[n][0] is the nth test image
# flower_test[n][1] is the nth test label
# flower_classes[k] is the name of the kth class
flower_train, flower_test, flower_classes = load_flower_data()

len(flower_train), len(flower_test)

(1020, 1020)

# Display a sample in Flowers 102 dataset
sample_idx = 0 # Choose an image index that you want to display
print("Label:", flower_classes[flower_train[sample_idx][1]])
flower_train[sample_idx][0]
```

A close-up photograph of several pink flowers with yellow centers, likely from a flowering plant. The petals are a soft pink color with darker pink veins radiating from the center. The centers are bright yellow with prominent stamens. The background is dark and out of focus.


```
working in indexes: https://pypi.org/simple/ https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting git+https://github.com/openai/CLIP.git
Cloning https://github.com/openai/CLIP.git to /tmp/pip-req-build-qsft5q7b
Running command git clone --filter=blob:none --quiet https://github.com/openai/CLIP.git /tmp/pip-req-build-qsft5q7b
Resolved https://github.com/openai/CLIP.git to commit a9b1bf5920416aaeac965c25dd9e8f98c864f16
Preparing metadata (setup.py) ... done
Collecting ftfy
  Downloading ftfy-6.1.1-py3-none-any.whl (53 kB)
53.1/53.1 KB 7.1 MB/s eta 0:00:00
Requirement already satisfied: regex in /usr/local/lib/python3.9/dist-packages (from clip==1.0) (2022.10.31)
Requirement already satisfied: tqdm in /usr/local/lib/python3.9/dist-packages (from clip==1.0) (4.65.0)
Requirement already satisfied: torch in /usr/local/lib/python3.9/dist-packages (from clip==1.0) (1.13.1+cu116)
Requirement already satisfied: torchvision in /usr/local/lib/python3.9/dist-packages (from clip==1.0) (0.14.1+cu116)
Requirement already satisfied: wcwidth>=0.2.5 in /usr/local/lib/python3.9/dist-packages (from ftfy->clip==1.0) (0.2.6)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.9/dist-packages (from torch->clip==1.0) (4.5.0)
Requirement already satisfied: requests in /usr/local/lib/python3.9/dist-packages (from torchvision->clip==1.0) (2.27.1)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.9/dist-packages (from torchvision->clip==1.0) (8.4.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.9/dist-packages (from torchvision->clip==1.0) (1.22.4)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.9/dist-packages (from requests->torchvision->clip==1.0) (1.26.15)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.9/dist-packages (from requests->torchvision->clip==1.0) (3.4)
Requirement already satisfied: charset-normalizer==2.0.0 in /usr/local/lib/python3.9/dist-packages (from requests->torchvision->clip==1.0) (2.0.12)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.9/dist-packages (from requests->torchvision->clip==1.0) (2022.12.7)
Building wheels for collected packages: clip
  Building wheel for clip (setup.py) ... done
  Created wheel for clip: filename=clip-1.0-py3-none-any.whl size=1369398 sha256=b91c43c2ffa7af7e98485d2396c2b6bcfd85bf87765f2462283d12b1b7bc0cce
  Stored in directory: /tmp/pip-ephem-wheel-cache-b8_j6m15/wheels/c8/e4/e1/11374c111837672fc2068dfbe0d4b424cb9cdd1b2e184a71b5
Successfully built clip
Installing collected packages: ftfy, clip
Successfully installed clip-1.0 ftfy-6.1.1
```

[illegible]

```
"""The following is an example of using CLIP pre-trained model for zero-shot prediction task"""
# Prepare the inputs
n = 200
image, class_id = flower_train[n]
image_input = clip_preprocess(image).unsqueeze(0).to(device) # extract image and put in device memory
text_inputs = torch.cat([clip.tokenize(f"a photo of a {c}, a type of flower.") for c in flower_classes]).to(device) # put text to match to image

# Calculate features
with torch.no_grad():
    image_features = clip_model.encode_image(image_input) # compute image features with CLIP model
    text_features = clip_model.encode_text(text_inputs) # compute text features with CLIP model
image_features /= image_features.norm(dim=-1, keepdim=True) # unit-normalize image features
text_features /= text_features.norm(dim=-1, keepdim=True) # unit-normalize text features

# Pick the top 5 most similar labels for the image
similarity = (100.0 * image_features @ text_features.T) # score is cosine similarity times 100
p_class_given_image = similarity.softmax(dim=-1) # P(y|x) is score through softmax
values, indices = p_class_given_image[0].topk(5) # gets the top 5 labels

# Print the probability of the top five labels
print("Ground truth:", flower_classes[class_id])
print("\nTop predictions:\n")
for value, index in zip(values, indices):
```

```
print(f"{flower_classes[index]:>16s):  {100 * value.item():.2f}%")
```

image

Ground truth: giant white arum lily

Top predictions:

```
giant white arum lily: 60.01%
lotus: 12.98%
siam tulip: 8.25%
anthurium: 3.55%
morning glory: 1.87%
```



▼ 1.4 YOUR TASK: Test CLIP zero-shot performance on Flowers 102

```
from tqdm import tqdm
from torch.utils.data import DataLoader
```

```
# Load flowers dataset again. This time, with clip_preprocess as transform
flower_train_trans, flower_test_trans, flower_classes = load_flower_data(img_transform=clip_preprocess)
```

```
def clip_zero_shot(data_set, classes):
    data_loader = DataLoader(data_set, batch_size=64, shuffle=False) # dataloader lets you process in batch which is way faster
    num_correct = 0
    num_iteration = len(data_set)
    for n in range(num_iteration):
        image, class_id = data_set[n]
        image_input = clip_preprocess(image).unsqueeze(0).to(device)
        text_inputs = torch.cat([clip.tokenize(f"a photo of a {c}, a type of flower.") for c in flower_classes]).to(device)
        with torch.no_grad():
            image_features = clip_model.encode_image(image_input) # compute image features with CLIP model
            text_features = clip_model.encode_text(text_inputs) # compute text features with CLIP model
            image_features /= image_features.norm(dim=-1, keepdim=True) # unit-normalize image features
            text_features /= text_features.norm(dim=-1, keepdim=True)
            similarity = (100.0 * image_features @ text_features.T) # score is cosine similarity times 100
            p_class_given_image = similarity.softmax(dim=-1) # P(y|x) is score through softmax
            values, indices = p_class_given_image[0].topk(5) # gets the top 5 labels
            if class_id == indices[0]:
                num_correct += 1
    accuracy = num_correct / num_iteration
    return accuracy
```

```
accuracy = clip_zero_shot(data_set=flower_test, classes=flower_classes)
print(f"\nAccuracy = {100*accuracy:.3f}%")
```

Accuracy = 67.843%

▼ 1.5 YOUR TASK: Test CLIP linear probe performance on Flowers 102

```

from sklearn.linear_model import LogisticRegression
from tqdm import tqdm
from torch.utils.data import DataLoader
import numpy as np

```

"""

In this part, train a linear classifier on CLIP features
 return: image features, labels in numpy format.

"""

```

def get_features(data_set):
    # Needs code here
    all_feature = []
    all_label = []
    with torch.no_grad():
        for images, labels in tqdm(DataLoader(data_set, batch_size = 100)):
            features = model.encode_image(images.to(device))
            all_feature.append(features)
            all_label.append(labels)
    return torch.cat(all_feature).cpu().numpy(), torch.cat(all_label).cpu().numpy()

```

```

# Calculate the image features
device = "cuda" if torch.cuda.is_available() else "cpu"
model, preprocess = clip.load('ViT-B/32', device)
train_features, train_labels = get_features(flower_train_trans)
test_features, test_labels = get_features(flower_test_trans)
classifier = LogisticRegression(random_state = 0, C=0.316, max_iter = 1000, verbose = 1)
classifier.fit(train_features, train_labels)
# Perform logistic regression
# Needs code here

```

```

# Evaluate using the logistic regression classifier
# Needs code here
predictions = classifier.predict(test_features)
accuracy = np.mean((test_labels == predictions).astype(float))
print(f"\nAccuracy = {100*accuracy:.3f}%")

```

```

100%|██████████████████| 11/11 [00:13<00:00, 1.25s/it]
100%|██████████████████| 11/11 [00:13<00:00, 1.24s/it]
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

```

```

Accuracy = 93.627%
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 3.8s finished

```

1.6 YOUR TASK: Evaluate a nearest-neighbor classifier on CLIP features

```

from scipy import stats
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

```

```

def knn(x_train, y_train, x_test, y_test, K=1):
    # Needs code here
    model = KNeighborsClassifier(n_neighbors = K)
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)
    accuracy = accuracy_score(y_test, y_pred)
    return accuracy

```

```

K = [1, 3, 5, 11, 21]
for n in K:
    accuracy = knn(train_features, train_labels, test_features, test_labels, K=n)
    print(f"\nAccuracy = {100*accuracy:.3f}%")

```

```

Accuracy = 84.118%

```

```

Accuracy = 83.137%

```

```

Accuracy = 84.608%

```

```

Accuracy = 85.000%

```

```

Accuracy = 79.804%

```

Part 2: Fine-Tune for Pets Image Classification

Include all the code for Part 2 in this section

2.1 Prepare Data

```

import torch
import torch.nn as nn
import torch.optim.lr_scheduler as lrs
from torch.utils.data import DataLoader
import torchvision
from torchvision import datasets
from torchvision import transforms
import matplotlib.pyplot as plt
from tqdm import tqdm

import os
from pathlib import Path
import numpy as np

# Mount and define data dir
from google.colab import drive
drive.mount('/content/drive')
datadir = "/content/"
save_dir = "/content/drive/My Drive/CS441/hw3"

def load_pet_dataset(train_transform = None, test_transform = None):
    OxfordIIITPet = datasets.OxfordIIITPet
    if os.path.isdir(datadir+ "oxford-iiit-pet"):
        do_download = False
    else:
        do_download = True
    training_set = OxfordIIITPet(root = datadir,
                                split = 'trainval',
                                transform = train_transform,
                                download = do_download)

    test_set = OxfordIIITPet(root = datadir,
                             split = 'test',
                             transform = test_transform,
                             download = do_download)

    return training_set, test_set

train_set, test_set = load_pet_dataset()

# Display a sample in OxfordIIIPet dataset
sample_idx = 0 # Choose an image index that you want to display
print("Label:", train_set.classes[train_set[sample_idx][1]])
train_set[sample_idx][0]

```

2.2 Data Preprocess

```

from torchvision import transforms
from torch.utils.data import DataLoader

# Feel free to add augmentation choices

# Apply data augmentation
train_transform = transforms.Compose([
    transforms.Resize(224),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                          std= [0.229, 0.224, 0.225]),
])

test_transform = transforms.Compose([
    transforms.Resize(224), # resize to 224x224 because that's the size of ImageNet images
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                          std= [0.229, 0.224, 0.225]),
])

# Feel free to change
train_set, test_set = load_pet_dataset(train_transform, test_transform)
train_loader = DataLoader(dataset=train_set,
                           batch_size=64,
                           shuffle=True,
                           num_workers=2)

test_loader = DataLoader(dataset=test_set,

```

```
batch_size=64,
shuffle=False,
num_workers=2)
```

▼ 2.3 Helper Functions

```
# Display the number of parameters and model structure
def display_model(model):
    # Check number of parameters
    summary_dict = {}
    num_params = 0
    summary_str = [' '*80]

    for module_name, module in model.named_children():
        summary_count = 0
        for name, param in module.named_parameters():
            if(param.requires_grad):
                summary_count += param.numel()
                num_params += param.numel()
        summary_dict[module_name] = [summary_count]
        summary_str+= [f'- {module_name: <40} : {str(summary_count):^34s}']

    summary_dict['total'] = [num_params]

    # print summary string
    summary_str += [' '*80]
    summary_str += ['--' + f'{"Total":<40} : {str(num_params) + " params":^34s}' + '--']
    print('\n'.join(summary_str))

    # print model structure
    print(model)

# Plot loss or accuracy
def plot_losses(train, val, test_frequency, num_epochs):
    plt.plot(train, label="train")
    indices = [i for i in range(num_epochs) if ((i+1)%test_frequency == 0 or i ==0 or i == 1)]
    plt.plot(indices, val, label="val")
    plt.title("Loss Plot")
    plt.ylabel("Loss")
    plt.xlabel("Epoch")
    plt.legend()
    plt.show()

def plot_accuracy(train, val, test_frequency, num_epochs):
    indices = [i for i in range(num_epochs) if ((i+1)%test_frequency == 0 or i ==0 or i == 1)]
    plt.plot(indices, train, label="train")
    plt.plot(indices, val, label="val")
    plt.title("Training Plot")
    plt.ylabel("Accuracy")
    plt.xlabel("Epoch")
    plt.legend()
    plt.show()

def save_checkpoint(save_dir, model, save_name = 'best_model.pth'):
    save_path = os.path.join(save_dir, save_name)
    torch.save(model.state_dict(), save_path)

def load_model(model, save_dir, save_name = 'best_model.pth'):
    save_path = os.path.join(save_dir, save_name)
    model.load_state_dict(torch.load(save_path))
    return model
```

▼ 2.4 YOUR TASK: Fine-Tune Pre-trained Network on Pets

Read and understand the code and then uncomment it. Then, set up your learning rate, learning scheduler, and train/evaluate. Adjust as necessary to reach target performance.

```
def train(train_loader, model, criterion, optimizer):
    """
    Train network
    :param train_loader: training dataloader
    :param model: model to be trained
    :param criterion: criterion used to calculate loss (should be CrossEntropyLoss from torch.nn)
    :param optimizer: optimizer for model's params (Adams or SGD)
    :return: mean training loss
    """

    model.train()
    loss_ = 0.0
```

```

losses = []

# TO DO: read this documentation and then uncomment the line below; https://pypi.org/project/tqdm/
# it_train = tqdm(enumerate(train_loader), total=len(train_loader), desc="Training ...", position = 0) # progress bar
for i, (images, labels) in it_train:

    # TO DO: read/understand and then uncomment these lines
    #images, labels = images.to(device), labels.to(device)
    #optimizer.zero_grad()
    #prediction = model(images)
    #loss = criterion(prediction, labels)
    #it_train.set_description(f'loss: {loss:.3f}')
    #loss.backward()
    #optimizer.step()
    #losses.append(loss)

return torch.stack(losses).mean().item()

def test(test_loader, model, criterion):
    """
    Test network.
    :param test_loader: testing dataloader
    :param model: model to be tested
    :param criterion: criterion used to calculate loss (should be CrossEntropyLoss from torch.nn)
    :return: mean_accuracy: mean accuracy of predicted labels
             test_loss: mean test loss during testing
    """
    model.eval()
    losses = []
    correct = 0
    total = 0

    # TO DO: read this documentation and then uncomment the line below; https://pypi.org/project/tqdm/
    #it_test = tqdm(enumerate(test_loader), total=len(test_loader), desc="Validating ...", position = 0)
    for i, (images, labels) in it_test:

        # TO DO: read/understand and then uncomment these lines
        #images, labels = images.to(device), labels.to(device)
        #with torch.no_grad(): # https://pytorch.org/docs/stable/generated/torch.no_grad.html
        #    output = model(images)
        #preds = torch.argmax(output, dim=-1)
        #loss = criterion(output, labels)
        #losses.append(loss.item())
        #correct += (preds == labels).sum().item()
        #total += len(labels)

    mean_accuracy = correct / total
    test_loss = np.mean(losses)
    print('Mean Accuracy: {:.4f}'.format(mean_accuracy))
    print('Avg loss: {}'.format(test_loss))

    return mean_accuracy, test_loss

device = 'cuda'
# loads a pre-trained ResNet-34 model
model = torch.hub.load('pytorch/vision:v0.10.0', 'resnet34', pretrained=True)
target_class = 37
# TO DO: replace the last layer with a new linear layer for Pets classification

model = model.to(device)
display_model(model) # displays the model structure and parameter count

# Training Setting. Feel free to change.
num_epochs = 20
test_interval = 5

# TO DO: set initial learning rate
learn_rate = []
optimizer = torch.optim.Adam(model.parameters(), lr=learn_rate)

# TO DO: define your learning rate scheduler, e.g. StepLR
# https://pytorch.org/docs/stable/generated/torch.optim.lr_scheduler.StepLR.html#torch.optim.lr_scheduler.StepLR
lr_scheduler = []

criterion = torch.nn.CrossEntropyLoss()

train_losses = []
train_accuracy_list = []
test_losses = []
test_accuracy_list = []

# Iterate over the DataLoader for training data

```

```

for epoch in tqdm(range(num_epochs), total=num_epochs, desc="Training ...", position=1):
    train_loss = train(train_loader, model, criterion, optimizer) # Train the Network for one epoch
    # TO DO: uncomment the line below. It should be called each epoch to apply the lr_scheduler
    # lr_scheduler.step()
    train_losses.append(train_loss)
    print(f'Loss for Training on epoch {str(epoch)} is {str(train_loss)} \n')

    if(epoch%test_interval==0 or epoch==1 or epoch==num_epochs-1):
        print('Evaluating Network')

        train_accuracy, _ = test(train_loader, model, criterion) # Get training accuracy
        train_accuracy_list.append(train_accuracy)

        print(f'Training accuracy on epoch {str(epoch)} is {str(train_accuracy)} \n')

        test_accuracy, test_loss = test(test_loader, model, criterion) # Get testing accuracy and error
        test_losses.append(test_loss)
        test_accuracy_list.append(test_accuracy)

        print(f'Testing accuracy on epoch {str(epoch)} is {str(test_accuracy)} \n')

    # Checkpoints are used to save the model with best validation accuracy
    if test_accuracy >= max(test_accuracy_list):
        print("Saving Model")
        save_checkpoint(save_dir, model, save_name = 'best_model.pth') # Save model with best performance

```

▼ 2.5 Plotting of losses and accuracy

```

plot_losses(train_losses, test_losses, test_interval, num_epochs)
plot_accuracy(train_accuracy_list, test_accuracy_list, test_interval, num_epochs)

```

▼ 2.6 Evaluating trained model

```

# TO DO: initialize your trained model as you did before so that you can load the parameters into it
model = torch.hub.load('pytorch/vision:v0.10.0', 'resnet34', pretrained=True).to(device)
# replace last layer

```

```

load_model(model, save_dir) # Load the trained weight

```

```

test_accuracy, test_loss= test(test_loader, model, criterion)
print(f"Testing accuracy is {str(test_accuracy)} \n")

```

Part 3: No coding for this part

▼ Part 4: Stretch Goals

Include any new code needed for Part 3 here

```

# example network definition that needs to be modified for custom network stretch goal

```

```

class Network(nn.Module):
    def __init__(self, num_classes=10, dropout = 0.5):
        super(Network, self).__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=11, stride=4, padding=2),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
            nn.Conv2d(64, 256, kernel_size=5, padding=2),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
            nn.Conv2d(256, 256, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
        )

        self.avgpool = nn.AdaptiveAvgPool2d((6, 6))
        self.classifier = nn.Sequential(
            nn.Dropout(p=dropout),
            nn.Linear(256 * 6 * 6, 512),
            nn.ReLU(inplace=True),
            nn.Dropout(p=dropout),
            nn.Linear(512, 512),
            nn.ReLU(inplace=True),
            nn.Linear(512, num_classes),
        )

```


)

```
def forward(self, x):
    N, c, H, W = x.shape
    features = self.features(x)
    pooled_features = self.avgpool(features)
    output = self.classifier(torch.flatten(pooled_features, 1))
    return output
```

✓ 0s completed at 7:04 PM

