

# CS441: Applied ML - HW 3

## ▼ Part 1: CLIP: Contrastive Language-Image Pretraining

Include all the code for Part 1 in this section

### ▼ 1.1 Prepare data

[Here](#) is the json file you need for labels of flowers 102

```
import json
import os
import os.path as osp
import numpy as np
from google.colab import drive
import torch
from torchvision.datasets import Flowers102
%matplotlib inline
from matplotlib import pyplot as plt
```

```
drive.mount('/content/drive')
datadir = "/content/drive/My Drive/CS441/hw3/"
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive")

```
def load_flower_data(img_transform=None):
    drive.mount('/content/drive')
    datadir = "/content/drive/My Drive/CS441/hw3/" #/content/drive/My Drive/CS441/hw3/"
    if os.path.isdir(datadir+ "flowers-102"):
        do_download = False
    else:
        do_download = True
    train_set = Flowers102(root=datadir, split='train', transform=img_transform, download=do_download)
    test_set = Flowers102(root=datadir, split='val', transform=img_transform, download=do_download)
    classes = json.load(open(osp.join("/content/drive/My Drive/CS441/hw3/", "flowers102_classes.json")))

    return train_set, test_set, classes
```

```
# READ ME! This takes some time (a few minutes), so if you are using Colabs,
# first set to use GPU: Edit->Notebook Settings->Hardware Accelerator=GPU, and r
```

```
# Data structure details
# flower_train[n][0] is the nth train image
# flower_train[n][1] is the nth train label
# flower_test[n][0] is the nth test image
# flower_test[n][1] is the nth test label
```

```
# flower_classes[k] is the name of the kth class
flower_train, flower_test, flower_classes = load_flower_data()
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive")

```
len(flower_train), len(flower_test)
```

```
(1020, 1020)
```

```
# Display a sample in Flowers 102 dataset
sample_idx = 0 # Choose an image index that you want to display
print("Label:", flower_classes[flower_train[sample_idx][1]])
flower_train[sample_idx][0]
```

Label: pink primrose



## ▼ 1.2 Prepare CLIP model

```
!pip install git+https://github.com/openai/CLIP.git
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting git+<https://github.com/openai/CLIP.git>

Cloning <https://github.com/openai/CLIP.git> to /tmp/pip-req-build-o5875vcp

Running command git clone --filter=blob:none --quiet <https://github.com/openai/CLIP.git> /tmp/pip-req-build-o5875vcp

Resolved <https://github.com/openai/CLIP.git> to commit a9b1bf5920416aaeac965c25dd9e8f98c864f16

Preparing metadata (setup.py) ... done

Collecting ftfy

53.1/53.1 KB 6.4 M

```
import clip
```

[illegible]

```

"""The following is an example of using CLIP pre-trained model for zero-shot prediction task"""
# Prepare the inputs
n = 200
image, class_id = flower_train[n]
image_input = clip_preprocess(image).unsqueeze(0).to(device) # extract image and put in device memory
text_inputs = torch.cat([clip.tokenize(f"a photo of a {c}, a type of flower.") for c in flower_classes])

# Calculate features
with torch.no_grad():
    image_features = clip_model.encode_image(image_input) # compute image features with CLIP model
    text_features = clip_model.encode_text(text_inputs) # compute text features with CLIP model
    image_features /= image_features.norm(dim=-1, keepdim=True) # unit-normalize image features
    text_features /= text_features.norm(dim=-1, keepdim=True) # unit-normalize text features

# Pick the top 5 most similar labels for the image
similarity = (100.0 * image_features @ text_features.T) # score is cosine similarity times 100
p_class_given_image = similarity.softmax(dim=-1) # P(y|x) is score through softmax

```

```
values, indices = p_class_given_image[0].topk(5) # gets the top 5 labels
```

```
# Print the probability of the top five labels
```

```
print("Ground truth:", flower_classes[class_id])
```

```
print("\nTop predictions:\n")
```

```
for value, index in zip(values, indices):
```

```
    print(f"{flower_classes[index]:>16s}: {100 * value.item():.2f}%")
```

```
image
```

Ground truth: giant white arum lily

Top predictions:

giant white arum lily: 60.59%

lotus: 12.80%

siam tulip: 8.07%

anthurium: 3.52%

morning glory: 1.85%



## ▼ 1.4 YOUR TASK: Test CLIP zero-shot performance on Flowers 102

```
from tqdm import tqdm
```

```
from torch.utils.data import DataLoader
```

```
# Load flowers dataset again. This time, with clip_preprocess as transform
```

```
flower_train_trans, flower_test_trans, flower_classes = load_flower_data(img_transform=clip_preprocess)
```

```
def clip_zero_shot(data_loader, classes):
```

```

def clip_zero_shot(data_set, classes):
    data_loader = DataLoader(data_set, batch_size=64, shuffle=False) # dataloader lets you
    num_correct = 0
    num_iteration = len(data_set)
    for n in range(num_iteration):
        image, class_id = data_set[n]
        image_input = clip_preprocess(image).unsqueeze(0).to(device)
        text_inputs = torch.cat([clip.tokenize(f"a photo of a {c}, a type of flower.") for c in classes], dim=0)
        with torch.no_grad():
            image_features = clip_model.encode_image(image_input) # compute image features with
            text_features = clip_model.encode_text(text_inputs) # compute text features with
            image_features /= image_features.norm(dim=-1, keepdim=True) # unit-normalize image features
            text_features /= text_features.norm(dim=-1, keepdim=True)
            similarity = (100.0 * image_features @ text_features.T) # score is cosine similarity
            p_class_given_image = similarity.softmax(dim=-1) # P(y|x) is score through softmax
            values, indices = p_class_given_image[0].topk(5) # gets the top 5 labels
            if class_id == indices[0]:
                num_correct += 1
    accuracy = num_correct / num_iteration
    return accuracy

```

```

accuracy = clip_zero_shot(data_set=flower_test, classes=flower_classes)
print(f"\nAccuracy = {100*accuracy:.3f}%")

```

Accuracy = 67.843%

## ▼ 1.5 YOUR TASK: Test CLIP linear probe performance on Flowers 102

```

from sklearn.linear_model import LogisticRegression
from tqdm import tqdm
from torch.utils.data import DataLoader
import numpy as np

```

"""

In this part, train a linear classifier on CLIP features  
 return: image features, labels in numpy format.

"""

```

def get_features(data_set):
    # Needs code here
    all_feature = []
    all_label = []
    with torch.no_grad():
        for images, labels in tqdm(DataLoader(data_set, batch_size = 100)):
            features = model.encode_image(images.to(device))
            all_feature.append(features)
            all_label.append(labels)
    return torch.cat(all_feature).cpu().numpy(), torch.cat(all_label).cpu().numpy()

```

```

# Calculate the image features
device = "cuda" if torch.cuda.is_available() else "cpu"
model, preprocess = clip.load('ViT-B/32', device)
train_features, train_labels = get_features(flower_train_trans)
test_features, test_labels = get_features(flower_test_trans)
classifier = LogisticRegression(random_state = 0, C=0.316, max_iter = 1000, verbose = 1)
classifier.fit(train_features, train_labels)
# Perform logistic regression
# Needs code here

# Evaluate using the logistic regression classifier
# Needs code here
predictions = classifier.predict(test_features)
accuracy = np.mean((test_labels == predictions).astype(float))
print(f"\nAccuracy = {100*accuracy:.3f}%")

100%|██████████| 11/11 [00:13<00:00, 1.25s/it]
100%|██████████| 11/11 [00:13<00:00, 1.24s/it]
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

Accuracy = 93.627%
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 3.8s finished

```

## ▼ 1.6 YOUR TASK: Evaluate a nearest-neighbor classifier on CLIP features

```

from scipy import stats
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

def knn(x_train, y_train, x_test, y_test, K=1):
    # Needs code here
    model = KNeighborsClassifier(n_neighbors = K)
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)
    accuracy = accuracy_score(y_test, y_pred)
    return accuracy

K = [1, 3, 5, 11, 21]
for n in K:
    accuracy = knn(train_features, train_labels, test_features, test_labels, K=n)
    print(f"\nAccuracy = {100*accuracy:.3f}%")

Accuracy = 84.118%

Accuracy = 83.137%

Accuracy = 84.608%

Accuracy = 85.000%

Accuracy = 79.804%

```

## ▼ Part 2: Fine-Tune for Pets Image Classification

Include all the code for Part 2 in this section

### ▼ 2.1 Prepare Data

```
import torch
import torch.nn as nn
import torch.optim.lr_scheduler as lrs
from torch.utils.data import DataLoader
import torchvision
from torchvision import datasets
from torchvision import transforms
import matplotlib.pyplot as plt
from tqdm import tqdm
```

```
import os
from pathlib import Path
import numpy as np
```

```
# Mount and define data dir
from google.colab import drive
drive.mount('/content/drive')
datadir = "/content/"
save_dir = "/content/drive/My Drive/CS441/hw3"
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive")

```
def load_pet_dataset(train_transform = None, test_transform = None):
    OxfordIIITPet = datasets.OxfordIIITPet
    if os.path.isdir(datadir+ "oxford-iiit-pet"):
        do_download = False
    else:
        do_download = True
    training_set = OxfordIIITPet(root = datadir,
                                  split = 'trainval',
                                  transform = train_transform,
                                  download = do_download)

    test_set = OxfordIIITPet(root = datadir,
                              split = 'test',
                              transform = test_transform,
                              download = do_download)

    return training_set, test_set
```

```
train_set, test_set = load_pet_dataset()
```

```
# Display a sample in OxfordIIIPet dataset
```



```
sample_idx = 0 # Choose an image index that you want to display
print("Label:", train_set.classes[train_set[sample_idx][1]])
train_set[sample_idx][0]
```

Label: Abyssinian



## ▼ 2.2 Data Preprocess

```
from torchvision import transforms
from torch.utils.data import DataLoader
```

```
# Feel free to add augmentation choices
```

```
# Apply data augmentation
```

```
train_transform = transforms.Compose([
    transforms.Resize(224),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                          std= [0.229, 0.224, 0.225]),
])
```

```
test_transform = transforms.Compose([
    transforms.Resize(224), # resize to 224x224 because that's the size of Image
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
```



```
std= [0.229, 0.224, 0.225]),
```

```
])
```

```
# Feel free to change
train_set, test_set = load_pet_dataset(train_transform, test_transform)
train_loader = DataLoader(dataset=train_set,
                           batch_size=64,
                           shuffle=True,
                           num_workers=2)

test_loader = DataLoader(dataset=test_set,
                          batch_size=64,
                          shuffle=False,
                          num_workers=2)
```

## ▼ 2.3 Helper Functions

```
# Display the number of parameters and model structure
def display_model(model):
    # Check number of parameters
    summary_dict = {}
    num_params = 0
    summary_str = [' '*80]

    for module_name, module in model.named_children():
        summary_count = 0
        for name, param in module.named_parameters():
            if param.requires_grad:
                summary_count += param.numel()
                num_params += param.numel()
        summary_dict[module_name] = [summary_count]
        summary_str+= [f'- {module_name: <40} : {str(summary_count):^34s}']

    summary_dict['total'] = [num_params]

    # print summary string
    summary_str += [' '*80]
    summary_str += ['--' + f'{"Total":<40} : {str(num_params) + " params":^34s}' + '--']
    print('\n'.join(summary_str))

    # print model structure
    print(model)

# Plot loss or accuracy
def plot_losses(train, val, test_frequency, num_epochs):
    plt.plot(train, label="train")
    indices = [i for i in range(num_epochs) if ((i+1)%test_frequency == 0 or i ==0 or i == num_epochs)]
    plt.plot(indices, val, label="val")
    plt.title("Loss Plot")
    plt.ylabel("Loss")
    plt.xlabel("Epoch")
```

```
plt.legend()
plt.show()
```

```
def plot_accuracy(train, val, test_frequency, num_epochs):
    indices = [i for i in range(num_epochs) if ((i+1)%test_frequency == 0 or i ==0 or i == num_epochs)]
    plt.plot(indices, train, label="train")
    plt.plot(indices, val, label="val")
    plt.title("Training Plot")
    plt.ylabel("Accuracy")
    plt.xlabel("Epoch")
    plt.legend()
    plt.show()

def save_checkpoint(save_dir, model, save_name = 'best_model.pth'):
    save_path = os.path.join(save_dir, save_name)
    torch.save(model.state_dict(), save_path)

def load_model(model, save_dir, save_name = 'best_model.pth'):
    save_path = os.path.join(save_dir, save_name)
    model.load_state_dict(torch.load(save_path))
    return model
```

## ▼ 2.4 YOUR TASK: Fine-Tune Pre-trained Network on Pets

Read and understand the code and then uncomment it. Then, set up your learning rate, learning scheduler, and train/evaluate. Adjust as necessary to reach target performance.

```
def train(train_loader, model, criterion, optimizer):
    """
    Train network
    :param train_loader: training dataloader
    :param model: model to be trained
    :param criterion: criterion used to calculate loss (should be CrossEntropyLoss from torch.nn)
    :param optimizer: optimizer for model's params (Adams or SGD)
    :return: mean training loss
    """
    model.train()
    loss_ = 0.0
    losses = []

    # TO DO: read this documentation and then uncomment the line below; https://pypi.org/project/tqdm/
    it_train = tqdm(enumerate(train_loader), total=len(train_loader), desc="Training ...", position=0)
    for i, (images, labels) in it_train:

        # TO DO: read/understand and then uncomment these lines
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
        prediction = model(images)
        loss = criterion(prediction, labels)
        it_train.set_description(f'loss: {loss:.3f}')
        loss.backward()
        optimizer.step()
        losses.append(loss)
```

```

        return torch.stack(losses).mean().item()

def test(test_loader, model, criterion):
    """
    Test network.
    :param test_loader: testing dataloader
    :param model: model to be tested
    :param criterion: criterion used to calculate loss (should be CrossEntropyLoss from torch.nn)
    :return: mean_accuracy: mean accuracy of predicted labels
             test_loss: mean test loss during testing
    """
    model.eval()
    losses = []
    correct = 0
    total = 0

    # TO DO: read this documentation and then uncomment the line below; https://pytorch.org/docs/stable/generated/torch.no_grad.html
    it_test = tqdm(enumerate(test_loader), total=len(test_loader), desc="Validating ...", position=0)
    for i, (images, labels) in it_test:

        # TO DO: read/understand and then uncomment these lines
        images, labels = images.to(device), labels.to(device)
        with torch.no_grad(): # https://pytorch.org/docs/stable/generated/torch.no_grad.html
            output = model(images)
            preds = torch.argmax(output, dim=-1)
            loss = criterion(output, labels)
            losses.append(loss.item())
            correct += (preds == labels).sum().item()
            total += len(labels)

    mean_accuracy = correct / total
    test_loss = np.mean(losses)
    print('Mean Accuracy: {:.4f}'.format(mean_accuracy))
    print('Avg loss: {}'.format(test_loss))

    return mean_accuracy, test_loss


device = 'cuda'
# loads a pre-trained ResNet-34 model
model = torch.hub.load('pytorch/vision:v0.10.0', 'resnet34', pretrained=True)
target_class = 37
# TO DO: replace the last layer with a new linear layer for Pets classification

model = model.to(device)
display_model(model) # displays the model structure and parameter count

```

Using cache found in /root/.cache/torch/hub/pytorch\_vision\_v0.10.0

```
RuntimeError                                Traceback (most recent call last)
<ipython-input-49-2db0690e2f8a> in <module>
      5 # TO DO: replace the last layer with a new linear layer for Pets classification
      6
----> 7 model = model.to(device)
      8 display_model(model) # displays the model structure and parameter count
```

⚡ 4 frames

```
/usr/local/lib/python3.9/dist-packages/torch/cuda/__init__.py in _lazy_init()
    227         if 'CUDA_MODULE_LOADING' not in os.environ:
```

```
# Training Setting. Feel free to change.
```

```
num_epochs = 20
```

```
test_interval = 5
```

```
# TO DO: set initial learning rate
```

```
learn_rate = 0.1
```

```
optimizer = torch.optim.Adam(model.parameters(), lr=learn_rate)
```

```
# TO DO: define your learning rate scheduler, e.g. StepLR
```

```
# https://pytorch.org/docs/stable/generated/torch.optim.lr\_scheduler.StepLR.html#torch.optim.lr\_scheduler.StepLR
```

```
lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer, gamma=0.1, step_size=20)
```

```
criterion = torch.nn.CrossEntropyLoss()
```

```
train_losses = []
```

```
train_accuracy_list = []
```

```
test_losses = []
```

```
test_accuracy_list = []
```

```
# Iterate over the DataLoader for training data
```

```
for epoch in tqdm(range(num_epochs), total=num_epochs, desc="Training ...", position=1):
```

```
    train_loss = train(train_loader, model, criterion, optimizer) # Train the Network for one
```

```
    # TO DO: uncomment the line below. It should be called each epoch to apply the lr_scheduler
```

```
    lr_scheduler.step()
```

```
    train_losses.append(train_loss)
```

```
    print(f'Loss for Training on epoch {str(epoch)} is {str(train_loss)} \n')
```

```
    if(epoch%test_interval==0 or epoch==1 or epoch==num_epochs-1):
```

```
        print('Evaluating Network')
```

```
        train_accuracy, _ = test(train_loader, model, criterion) # Get training accuracy
```

```
        train_accuracy_list.append(train_accuracy)
```

```
        print(f'Training accuracy on epoch {str(epoch)} is {str(train_accuracy)} \n')
```

```
        test_accuracy, test_loss = test(test_loader, model, criterion) # Get testing accuracy
```

```
        test_losses.append(test_loss)
```

```
        test_accuracy_list.append(test_accuracy)
```

```
        print(f'Testing accuracy on epoch {str(epoch)} is {str(test_accuracy)} \n')
```

```
# Checkpoints are used to save the model with best validation accuracy
```

```
if test_accuracy >= max(test_accuracy_list):
```

```
print("Saving Model")
save_checkpoint(save_dir, model, save_name = 'best_model.pth') # Save model with
```

## 2.5 Plotting of losses and accuracy

```
plot_losses(train_losses, test_losses, test_interval, num_epochs)
plot_accuracy(train_accuracy_list, test_accuracy_list, test_interval, num_epochs)
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-35-e4a4885c23bc> in <module>
----> 1 plot_losses(train_losses, test_losses, test_interval, num_epochs)
      2 plot_accuracy(train_accuracy_list, test_accuracy_list, test_interval, num_epochs)

NameError: name 'train_losses' is not defined
```

SEARCH STACK OVERFLOW

## 2.6 Evaluating trained model

```
# • TO • DO: • initialize • your • trained • model • as • you • did • before • so • that • you • can • load • the • para
model = • torch.hub.load('pytorch/vision:v0.10.0', 'resnet34', pretrained=True).to(device) •
# • replace • last • layer
target_class = • 37
num_features = • model.fc.in_features
model.fc = • nn.Linear(num_features, target_class)
model = • model.to(device)
load_model(model, • save_dir) • # • Load • the • trained • weight

test_accuracy, • test_loss= • test(test_loader, • model, • criterion)
print(f"Testing • accuracy • is • {str(test_accuracy)} • \n")
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-2-0df689be9acf> in <module>
      1 # TO DO: initialize your trained model as you did before so that you can load the parameters
into it
----> 2 model = torch.hub.load('pytorch/vision:v0.10.0', 'resnet34', pretrained=True).to(device)
      3 # replace last layer
      4 target_class = 37
      5 num_features = model.fc.in_features

NameError: name 'torch' is not defined
```

SEARCH STACK OVERFLOW

## Part 3: No coding for this part

## Part 4: Stretch Goals



## Include any new code needed for Part 3 here

# example network definition that needs to be modified for custom network stretch goal

```
class Network(nn.Module):
    def __init__(self, num_classes=10, dropout = 0.5):
        super(Network, self).__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=11, stride=4, padding=2),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
            nn.Conv2d(64, 256, kernel_size=5, padding=2),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
            nn.Conv2d(256, 256, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
        )

        self.avgpool = nn.AdaptiveAvgPool2d((6, 6))
        self.classifier = nn.Sequential(
            nn.Dropout(p=dropout),
            nn.Linear(256 * 6 * 6, 512),
            nn.ReLU(inplace=True),
            nn.Dropout(p=dropout),
            nn.Linear(512, 512),
            nn.ReLU(inplace=True),
            nn.Linear(512, num_classes),
        )

    def forward(self, x):
        N, c, H, W = x.shape
        features = self.features(x)
        pooled_features = self.avgpool(features)
        output = self.classifier(torch.flatten(pooled_features, 1))
        return output
```

